

Operating system & Computer Networks Laboratory

Program	Master of Computer Applications
Semester	1
Course Title	Operating system & Computer Networks Laboratory
Course Code	23OMC109
Course Credits	2
Course Type	Core Course

1. Course Summary

The Operating System and Computer Network Lab course offers practical experience in computer networks and operating systems. Students learn network commands, configure network devices, implement point-to-point networks, and analyze network performance. They also simulate CPU scheduling algorithms, memory management techniques, disk scheduling algorithms. This hands-on approach enhances their understanding of network protocols, resource allocation, and system performance. By completing the term works, students gain valuable skills in network configuration, command-line tools, and practical implementation of operating system algorithms..

2. Course Outcomes (COs)

After the successful completion of this course, the student will be able to:

- CO-1.** Explain the use of network commands in command prompt to diagnose and troubleshoot network connectivity issues. (L2)
- CO-2.** Configure and demonstrate the use of hub, switch, and router in a simple network topology. (L3)
- CO-3.** Analyze the impact of queue size and bandwidth on network performance by implementing a point-to-point network with varying parameters and observing packet drops. (L4)
- CO-4.** Simulate CPU scheduling algorithms such as FCFS and SJF and analyze their performance using different process arrival times and burst times. (L4)
- CO-5.** Use the semaphore mechanism to control access to the shared buffer, ensuring proper synchronization between producer and consumer threads. (L4)

Operating systems

1. Simulate the following CPU scheduling algorithms-FCFS, SJF
2. Simulate the following Memory management algorithm-First fit, Best fit
3. Implement the optimal page replacement algorithm
4. Implement the FCFS Disk scheduling and SSTF Disk scheduling
5. Implement the producer consumer problem with solution using semaphore
6. Implement a program in C to extract process ID (PID) and parent process ID (PPID).

Experiment 1: CPU Scheduling Algorithms - FCFS and SJF

Objective: To simulate the First-Come-First-Serve (FCFS) and Shortest Job First (SJF) CPU scheduling algorithms.

Procedure:

FCFS Scheduling:

Write a C program that simulates the FCFS scheduling algorithm.

Input a list of processes with their burst times.

Calculate the waiting time and turnaround time for each process.

Display the scheduling order, waiting time, and turnaround time for each process.

SJF Scheduling:

Write a C program that simulates the SJF scheduling algorithm.

Input a list of processes with their burst times.

Sort the processes based on burst time (shortest job first).

Calculate the waiting time and turnaround time for each process.

Display the scheduling order, waiting time, and turnaround time for each process.

```
/* FCFS CPU Scheduling */
#include <stdio.h>
#include <stdlib.h>
struct Process {
    int pid; // process ID
    int bt; // burst time
    int at; // arrival time
};
void fcfs_scheduling(struct Process *proc, int n) {
    int i, wt[n], tat[n], total_wt = 0, total_tat = 0;
    // calculate waiting time for each process
    wt[0] = 0;
    for (i = 1; i < n; i++) {
```

```

    wt[i] = wt[i-1] + proc[i-1].bt;
}
// calculate turnaround time for each process
for (i = 0; i < n; i++) {
    tat[i] = wt[i] + proc[i].bt;
}
// calculate total waiting and turnaround time
for (i = 0; i < n; i++) {
    total_wt += wt[i];
    total_tat += tat[i];
}
// print result
printf("\nPID\tBT\tAT\tWT\tTAT\n");
for (i = 0; i < n; i++) {
    printf("%d\t%d\t%d\t%d\t%d\n", proc[i].pid, proc[i].bt, proc[i].at, wt[i], tat[i]);
}
printf("\nAverage waiting time: %.2f\n", (float)total_wt / n);
printf("Average turnaround time: %.2f\n", (float)total_tat / n);
}

int main() {
    int n, i;

    printf("Enter the number of processes: ");
    scanf("%d", &n);
    struct Process proc[n];
    for (i = 0; i < n; i++) {
        printf("Enter the burst time and arrival time for process %d: ", i+1);
        scanf("%d%d", &proc[i].bt, &proc[i].at);
        proc[i].pid = i+1;
    }
}

```

```
    fcfs_scheduling(proc, n);  
    return 0;  
}
```

SAMPLE OUTPUT

Enter the number of processes: 4

Enter the burst time and arrival time for process 1: 3 0

Enter the burst time and arrival time for process 2: 3 0

Enter the burst time and arrival time for process 3: 2 0

Enter the burst time and arrival time for process 4: 5 0

PID	BT	AT	WT	TAT
1	3	0	0	3
2	3	0	3	6
3	2	0	6	8
4	5	0	8	13

Average waiting time: 4.25

Average turnaround time: 7.50

```
/* SJF Final */  
#include <stdio.h>  
#include <stdlib.h>  
struct Process{  
    int pid;  
    int burst_time;  
    int waiting_time;  
    int turnaround_time;  
};  
void main()
```

```

{
    int n, i, j;
    struct Process temp;
    float avg_waiting_time = 0, avg_turnaround_time = 0;
    struct Process *processes;
    printf("Enter the number of processes: ");
    scanf("%d", &n);
    // Allocate memory for the array of processes
    processes = (struct Process*) malloc (n * sizeof (struct Process) );

    // Reading burst time of each process
    for(i=0; i<n; i++) {
        printf("Enter the burst time for process %d: ", i+1);
        scanf("%d", &processes[i].burst_time);
        processes[i].pid = i+1;
    }
    // Sort the processes based on their burst time in ascending order
    for(i=0; i<n-1; i++) {
        for(j=i+1; j<n; j++) {
            if(processes[i].burst_time > processes[j].burst_time) {
                // Swap the processes
                temp = processes[i];
                processes[i] = processes[j];
                processes[j] = temp;
            }
        }
    }
    // Calculate waiting time and turnaround time for each process
    processes[0].waiting_time = 0;
    processes[0].turnaround_time = processes[0].burst_time;
    for(i=1; i<n; i++) {

```

```

    processes[i].waiting_time = processes[i-1].waiting_time + processes[i-1].burst_time;
    processes[i].turnaround_time = processes[i].waiting_time + processes[i].burst_time;
}
// Calculate average waiting time and turnaround time
for(i=0; i<n; i++) {
    avg_waiting_time += processes[i].waiting_time;
    avg_turnaround_time += processes[i].turnaround_time;
}
avg_waiting_time /= n;
avg_turnaround_time /= n;
// Print the results
printf("Process ID\tBurst Time\tWaiting Time\tTurnaround Time\n");
// Sort the processes based on their ID in ascending order to print in proper order
for(i=0; i<n-1; i++) {
    for(j=i+1; j<n; j++) {
        if(processes[i].pid > processes[j].pid) {
            // Swap the processes
            temp = processes[i];
            processes[i] = processes[j];
            processes[j] = temp;
        }
    }
}

for(i=0; i<n; i++) {
    printf("%d\t%d\t%d\t%d\n", processes[i].pid, processes[i].burst_time, processes[i].waiting_time,
processes[i].turnaround_time);
}
printf("\nAverage Waiting Time = %0.2f ms", avg_waiting_time);
printf("\nAverage Turnaround Time = %0.2f ms\n", avg_turnaround_time);
// Free the allocated memory

```

```
    free(processes);  
}
```

In this program, we define a structure Process to store information about each process, including its process ID, burst time, waiting time, and turnaround time. The program first reads the burst time of each process and then sorts them in ascending order based on their burst time. It then calculates the waiting time and turnaround time for each process and finally prints the results, including the average waiting time and turnaround time.

SAMPLE OUTPUT 1:

Enter the number of processes: 3

Enter the burst time for process 1: 5

Enter the burst time for process 2: 2

Enter the burst time for process 3: 8

Process ID	Burst Time	Waiting Time	Turnaround Time
1	5	2	7
2	2	0	2
3	8	7	15

Average Waiting Time = 3.000000

Average Turnaround Time = 8.000000

SAMPLE OUTPUT 2:

Enter the number of processes: 4

Enter the burst time for process 1: 6

Enter the burst time for process 2: 8

Enter the burst time for process 3: 5

Enter the burst time for process 4: 2

Process ID	Burst Time	Waiting Time	Turnaround Time
1	6	7	13
2	8	13	21
3	5	2	7
4	2	0	2

Average Waiting Time = 5.500000

Average Turnaround Time = 10.750000

Experiment 2: Memory Management Algorithms - First Fit and Best Fit

Objective: To simulate the First Fit and Best Fit memory allocation algorithms.

Procedure:

First Fit Allocation:

Write a C program that simulates the First Fit memory allocation algorithm.

Maintain a list of available memory blocks and their sizes.

Input a list of processes with their memory requirements.

Allocate memory to processes based on the First Fit algorithm.

Display the memory allocation details.

Best Fit Allocation:

Write a C program that simulates the Best Fit memory allocation algorithm.

Maintain a list of available memory blocks and their sizes.

Input a list of processes with their memory requirements.

Allocate memory to processes based on the Best Fit algorithm.

Display the memory allocation details.

// C program to implement First Fit memory allocation algorithm

```
#include<stdio.h>

// Function to allocate memory
void firstFit(int m, int n, int Blocks[], int Process[])
{
    int i, j;
    // Stores block id of the
    int allocation[n]; // To store the allocated block number
    // Initialize to -1 as no block is assigned to processes
    for(i = 0; i < n; i++)
    {
        allocation[i] = -1;
```

```

}
// Look for available block to the process and allocate if fits

for (i = 0; i < n; i++)    //n number of processes
{
    for (j = 0; j < m; j++)    //m number of blocks
    {
        if (Blocks[j] >= Process[i])
        {
            // allocating block j to the ith process
            allocation[i] = j;

            // Reduce available memory in this block.
            Blocks[j] = Blocks[j] - Process[i];

            break; //go to the next process in the queue
        }
    }
}

printf("\nProcess No.\tProcess Size\t\tBlock no.\n");
for (int i = 0; i < n; i++)
{
    printf(" %d\t\t", i+1);
    printf("%d\t\t", Process[i]);
    if (allocation[i] != -1)
        printf("%i", allocation[i] + 1);
    else
        printf("Not Allocated");
}

```

```

        printf("\n");
    }
}

int main()
{
    int m; //number of blocks in the memory
    int n; //number of processes in the input queue
    int i;
    int Blocks[10];
    int Process[10];
    printf("Enter the number of blocks\n");
    scanf("%d", &m);
    printf("Enter the number of processes\n");
    scanf("%d", &n);
    printf("Enter the block sizes\n");
    for(i=0; i<m; i++)
        scanf("%d", &Blocks[i]);
    printf("Enter the process sizes\n");
    for(i=0; i<n; i++)
        scanf("%d", &Process[i]);

    firstFit(m, n, Blocks, Process);
    return 0 ;
}

```

SAMPLE OUTPUT 1

Enter the number of blocks

4

Enter the number of processes

3

Enter the block sizes

220

300

200

300

Enter the process sizes

250

100

600

Process No.	Process Size	Block no.
1	250	2
2	100	1
3	600	Not Allocated

SAMPLE OUTPUT 2

Enter the number of blocks

5

Enter the number of processes

4

Enter the block sizes

100

500

200

300

600

Enter the process sizes

225

450

120

350

Process No.	Process Size	Block no.
1	225	2
2	450	5
3	120	2
4	350	Not Allocated

SAMPLE OUTPUT 3

Enter the number of blocks

5

Enter the number of processes

4

Enter the block sizes

100

500

200

300

600

Enter the process sizes

220

430

110

425

Process No.	Process Size	Block no.
1	220	2
2	430	5
3	110	2
4	425	Not Allocated

```
// Implementation of Best Fit memory allocation algorithm in C
#include<stdio.h>
// Function to allocate memory
void BestFit(int m, int n, int partitions[], int process[])
{
    int i, j;
    int allocation[n]; // To store the partition allocated to process
    // Initialize to -1 as no block is assigned to processes
    for(i = 0; i < n; i++)
    {
        allocation[i] = -1;
    }
    // For every process select the best fit partition
    for (int i = 0; i < n; i++) // Select process
    {
        int bestPart = -1;
        for (int j = 0; j < m; j++) //Search the partition
        {
            if (partitions[j] >= process[i])
            {
                if (bestPart == -1)
                    bestPart = j;
            }
        }
    }
}
```

```

        else if (partitions[bestPart] > partitions[j])
            bestPart = j;
    }
}

// If bestPart is found
if (bestPart != -1)
{
    // allocate Partition j to p[i] process
    allocation[i] = bestPart;
    // Reduce the partition size of this block
    partitions[bestPart] = partitions[bestPart] - process[i];
}
}

printf("\nProcess No.\tProcess Size\t\tBlock no.\n");
for (int i = 0; i < n; i++)
{
    printf(" %d\t\t\t", i+1);
    printf("%d\t\t\t", process[i]);
    if (allocation[i] != -1)
        printf("%d", allocation[i] + 1);
    else
        printf("Not Allocated");
    printf("\n");
}
}

int main()
{

```

```

int m; //number of partitions in the memory
int n; //number of processes in the input queue
int i;
int partitions[10];
int process[10];
printf("Enter the number of partitions\n");
scanf("%d", &m);
printf("Enter the number of processes\n");
scanf("%d", &n);
printf("Enter the partition sizes\n");
for(i=0; i<m; i++)
    scanf("%d", &partitions[i]);
printf("Enter the process sizes\n");
for(i=0; i<n; i++)
    scanf("%d", &process[i]);
BestFit(m, n, partitions, process);
return 0 ;
}

```

SAMPLE OUTPUT 1

Enter the number of partitions

4

Enter the number of processes

3

Enter the partition sizes

220

300

200

300

Enter the process sizes

250

100

600

Process No.	Process Size	Block no.
1	250	2
2	100	3
3	600	Not Allocated

SAMPLE OUTPUT 2

Enter the number of partitions

5

Enter the number of processes

4

Enter the partition sizes

100

500

200

300

600

Enter the process sizes

225

650

120

350

Process No.	Process Size	Block no.
-------------	--------------	-----------

1	225	4
2	650	Not Allocated
3	120	3
4	350	2

SAMPLE OUTPUT 3

Enter the number of partitions

5

Enter the number of processes

4

Enter the partition sizes

100

500

200

300

600

Enter the process sizes

220

430

110

425

Process No.	Process Size	Block no.
1	220	4
2	430	2
3	110	3
4	425	5

SAMPLE OUTPUT 4

Enter the number of partitions

3

Enter the number of processes

2

Enter the partition sizes

100

600

300

Enter the process sizes

350

200

Process No.	Process Size	Block no.
1	350	2
2	200	2

Experiment 3: Optimal Page Replacement Algorithm

Objective: To implement and evaluate the Optimal Page Replacement algorithm for virtual memory management.

Procedure:

Write a C program to simulate the Optimal Page Replacement algorithm.

Input the reference string and the number of frames in memory.

Implement the Optimal algorithm to replace pages when a page fault occurs.

Calculate and display the page fault rate and the pages replaced.

```
/* Program to implement Optimal page replacement algorithm */
#include<stdio.h>
int main()
{
    int num_frames, num_pages, frames[10], pages[30], temp[10];
    int flag1, flag2, flag3, i, j, k, pos, max, faults = 0, hit;
    printf("Enter number of frames: ");
    scanf("%d", &num_frames);

    printf("Enter number of pages: ");
    scanf("%d", &num_pages);

    printf("Enter page reference string: ");
    for(i = 0; i < num_pages; i++){
        scanf("%d", &pages[i]);
    }

    /* assigning initial value */
    for(i = 0; i < num_frames; i++){
        frames[i] = -1;
    }
}
```

```

for(i = 0; i < num_pages; i++) { // Trying to pagein pages[i]
    flag1 = flag2 = 0;
    hit = 0;
    for(j = 0; j < num_frames; j++) {
        if(frames[j] == pages[i]){ // Page hit
            flag1 = flag2 = 1;
            hit = 1;
            break;
        }
    }
}

if(flag1 == 0){ // Page fault
    for(j = 0; j < num_frames; j++){ // Trying to find free frame
        if(frames[j] == -1){
            faults++;
            frames[j] = pages[i]; // pagein pages[i]
            flag2 = 1;
            break;
        }
    }
}

if(flag2 == 0){ // No free frame
    flag3 = 0;
    for(j = 0; j < num_frames; j++){ // Trying to find victim page
        temp[j] = -1;
        for(k = i + 1; k < num_pages; k++){ // Consider future reference
            if(frames[j] == pages[k]){ // pages[k] will be referred

```

```
temp[j] = k;
```

```
break;
```

```
}
```

```
}
```

```
}
```

```
for(j = 0; j < num_frames; j++){
```

```
if(temp[j] == -1){ // if temp[j] == -1 then that page will not be referred
```

```
pos = j; // use this pos to page in
```

```
flag3 = 1;
```

```
break;
```

```
}
```

```
}
```

```
if(flag3 == 0){ // All the pages in the frame will be referred
```

```
max = temp[0];
```

```
pos = 0;
```

```
for(j = 1; j < num_frames; j++){ // Find the page that will not be referred for long time
```

```
if(temp[j] > max){
```

```
max = temp[j];
```

```
pos = j; // use this pos to page in
```

```
}
```

```
}
```

```
}
```

```
frames[pos] = pages[i]; // Page in
```

```
faults++;
```

```
}
```

```
if (hit == 0)
```

```
{
```

```

        printf("\n");

        for(j = 0; j < num_frames; j++){
            printf("%d\t", frames[j]);
        }
    }
}
printf("\n\nTotal Page Faults = %d", faults);
return 0;
}

```

SAMPLE OUTPUT

Enter number of frames: 3

Enter number of pages: 10

Enter page reference string: 2

3

4

2

1

3

7

5

4

3

2	-1	-1
---	----	----

2	3	-1
---	---	----

2	3	4
---	---	---

1	3	4
---	---	---

7	3	4
5	3	4

Total Page Faults = 6

Experiment 4: Disk Scheduling - FCFS and SSTF

Objective: To implement and compare the First-Come-First-Serve (FCFS) and Shortest Seek Time First (SSTF) disk scheduling algorithms.

Procedure:

FCFS Disk Scheduling:

Write a C program to simulate the FCFS disk scheduling algorithm.

Input the sequence of disk track requests.

Calculate and display the total head movement.

SSTF Disk Scheduling:

Write a C program to simulate the SSTF disk scheduling algorithm.

Input the sequence of disk track requests.

Implement the SSTF algorithm to find the next track to move to.

Calculate and display the total head movement.

```
// C program to implement FCFS Disk Scheduling algorithm
#include <stdio.h>
void FCFS(int cylinder[],int num_cyli, int head_pos)
{
    int movement = 0;
    int cur_cyli, distance;
    printf("\nHead Movements\n");
    printf("\nFrom\tTo\tMovement\n");

    for( int i = 0 ; i < num_cyli ; i++ )
    {
        cur_cyli = cylinder[i];

        // Find absolute distance
        distance = fabs(head_pos - cur_cyli);
```

```

        //Increment the total count
        movement = movement + distance;

        printf("%d\t%d\t%d\n" , head_pos, cur_cyli, distance);

        //Assign new head position
        head_pos = cur_cyli;

    }
    printf(" Total number of head movemnts = %d ", movement);
}
int main()
{
    int cylinder[20]; // request array
    int num_cyli, i, head_pos;

    printf("Enter the number of requests\n");
    scanf("%d", &num_cyli);

    printf("Enter the disk queue of cylinders\n");
    for(i = 0 ; i < num_cyli ; i++)
        scanf("%d", &cylinder[i]);

    printf("Enter the initial head position\n");
    scanf("%d", &head_pos);
    FCFS(cylinder, num_cyli, head_pos);

    return 0;
}

```

SAMPLE OUTPUT

Enter the number of requests

8

Enter the disk queue of cylinders

98

183

37

122

14

124

65

67

Enter the initial head position

53

Head Movements

From	To	Movement
53	98	45
98	183	85
183	37	146
37	122	85
122	14	108
14	124	110
124	65	59
65	67	2

Total number of head movemnts = 640

```
//To implement the Shortest Seek Time First disk scheduling algorithm
#include<math.h>
```

```

#include<stdio.h>

void sjfs(int cylinder[], int num_cyli, int head_pos)
{
    int i,k,movement = 0,cur_pos,index[20],minimum,temp[20],j=0,min_index,distance;
    cur_pos = head_pos;
    for(k = 0 ; k < num_cyli ; k++)
    {
        for(i = 0 ; i < num_cyli ; i++)
        {
            // Calculate distance of each request from current position
            index[i]=fabs(cur_pos - cylinder[i]);
        }
        // Find the nearest request
        minimum = index[0];
        min_index = 0;
        for(i = 1 ; i < num_cyli ; i++)
        {
            if(minimum > index[i])
            {
                minimum = index[i];
                min_index = i;
            }
        }
        temp[j] = cylinder[min_index];
        j++;
        // change the current position value to next request
        cur_pos = cylinder[min_index];
        // Change the processed request so that it is not processed again
        cylinder[min_index] = 9999;
    }
}

```

```

}

printf("\nHead Movements\n");
    printf("\nFrom\tTo\tMovement\n");

    movement = movement + fabs(head_pos - temp[0]); // head movement
    printf("%d\t%d\t%d\n", head_pos, temp[0], movement);
    for(i = 1 ; i < num_cyli ; i++)
    {
        distance = fabs(temp[i]-temp[i-1]); ///head movement

        printf("%d\t%d\t%d\n", temp[i-1], temp[i], distance);
        movement = movement + distance;

    }
    printf("\n");
    printf("Total number of head movements = %d\n",movement);
}

int main()
{
    int cylinder[20]; // request array
    int num_cyli, i, head_pos;
    printf("Enter the initial head position\n");
    scanf("%d", &head_pos);

    printf("Enter the number of requests\n");
    scanf("%d", &num_cyli);

    printf("Enter the disk queue of cylinders\n");
    for(i = 0 ; i < num_cyli ; i++)

```

```
        scanf("%d", &cylinder[i]);  
        sjfs(cylinder, num_cyli, head_pos);  
        return 0;  
}
```

/* SAMPLE OUTPUT 1 */

Enter the initial head position

53

Enter the number of requests

8

Enter the disk queue of cylinders

98

183

37

122

14

124

65

67

Head Movements

From	To	Movement
------	----	----------

53	65	12
----	----	----

65	67	2
----	----	---

67	37	30
----	----	----

37	14	23
----	----	----

14	98	84
----	----	----

98	122	24
----	-----	----

122	124	2
-----	-----	---

124 183 59

Total number of head movements = 236

/* SAMPLE OUTPUT 2 */

Enter the initial head position

50

Enter the number of requests

2

Enter the disk queue of cylinders

60

10

Head Movements

From	To	Movement
------	----	----------

50	60	10
----	----	----

60	10	50
----	----	----

Total number of head movements = 60

Experiment 5: Producer-Consumer Problem with Semaphores

Objective: To implement a solution to the Producer-Consumer problem using semaphores in C.

Procedure:

Write a C program that demonstrates the Producer-Consumer problem.

Implement the producer and consumer functions using semaphores for synchronization.

Ensure that producers and consumers work concurrently without issues.

```
// C program for the Producer-Consumer
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
// Initialize a mutex to 1
```

```
int mutex = 1;
```

```
// Number of full slots as 0
```

```
int full = 0;
```

```
// Number of empty slots as size
```

```
// of buffer
```

```
int empty = 10, x = 0;
```

```
// Function to produce an item and
```

```
// add it to the buffer
```

```
void producer()
```

```
{
```

```
    // Decrease mutex value by 1
```

```
    --mutex;
```

```
    // Increase the number of full
```



```

        // slots by 1
        ++full;

        // Decrease the number of empty
        // slots by 1
        --empty;

        // Item produced
        x++;
        printf("\nProducer produces"
               "item %d",
               x);

        // Increase mutex value by 1
        ++mutex;
    }

// Function to consume an item and
// remove it from buffer
void consumer()
{
    // Decrease mutex value by 1
    --mutex;

    // Decrease the number of full
    // slots by 1
    --full;

    // Increase the number of empty
    // slots by 1

```

```

        ++empty;

        printf("\nConsumer consumes "
               "item %d",
               x);

        x--;

        // Increase mutex value by 1
        ++mutex;
    }

// Driver Code
int main()
{
    int n, i;

    printf("\n1. Press 1 for Producer"
           "\n2. Press 2 for Consumer"
           "\n3. Press 3 for Exit");

    // Using '#pragma omp parallel for'
    // can give wrong value due to
    // synchronization issues.

    // 'critical' specifies that code is
    // executed by only one thread at a
    // time i.e., only one thread enters
    // the critical section at a given time
    #pragma omp critical

    for (i = 1; i > 0; i++) {

```

```
printf("\nEnter your choice:");  
scanf("%d", &n);
```

```
// Switch Cases
```

```
switch (n) {
```

```
case 1:
```

```
    // If mutex is 1 and empty
```

```
    // is non-zero, then it is
```

```
    // possible to produce
```

```
    if ((mutex == 1)
```

```
        && (empty != 0)) {
```

```
        producer();
```

```
    }
```

```
    // Otherwise, print buffer
```

```
    // is full
```

```
    else {
```

```
        printf("Buffer is full!");
```

```
    }
```

```
    break;
```

```
case 2:
```

```
    // If mutex is 1 and full
```

```
    // is non-zero, then it is
```

```
    // possible to consume
```

```
    if ((mutex == 1)
```

```
        && (full != 0)) {
```

```
        consumer();
```

```

        }

        // Otherwise, print Buffer
        // is empty
        else{
            printf("Buffer is empty!");
        }
        break;

    // Exit Condition
    case 3:
        exit(0);
        break;
    }
}
}

```

SAMPLE OUTPUT:

1. Press 1 for Producer
2. Press 2 for Consumer
3. Press 3 for Exit

Enter your choice:1

Producer produces item 1

Enter your choice:1

Producer produces item 2

Enter your choice:1

Producer produces item 3

Enter your choice:2

Consumer consumes item 3

Enter your choice:2

Consumer consumes item 2

Enter your choice:2

Consumer consumes item 1

Enter your choice:2

Buffer is empty!

Enter your choice:1

Producer produces item 1

Enter your choice:2

Consumer consumes item 1

Enter your choice:2

Buffer is empty!

Enter your choice:3

...Program finished with exit code 0

Experiment 6: Process ID and Parent Process ID Extraction

Objective: To extract the Process ID (PID) and Parent Process ID (PPID) of a process in a C program.

Procedure:

Write a C program to obtain the Process ID (PID) and Parent Process ID (PPID) using system calls.

Display the PID and PPID of the current process.

Explain the significance of PID and PPID in process management.

```
#include<stdio.h>
#include<unistd.h>
#include<stdlib.h>
int main( )
{
int pid;
pid=fork( );
if(pid== -1)
{
perror("fork failed");
exit(0);
}
if(pid==0)
{
printf("\n1-1 Child process is under execution");
printf("\n1-2 Process id of the child process is %d", getpid());
printf("\n1-3 Process id of the parent process is %d",getppid());
}
else
{
printf("\n2-1 Parent process is under execution");
printf("\n2-2 Process id of the parent process is %d", getpid());
printf("\n2-3 Process id of the parent of parent is %d", getppid());
```

```
}  
return(0);  
}
```

SAMPLE OUTPUT:

2-1 Parent process is under execution

2-2 Process id of the parent process is 1861

2-3 Process id of the parent of parent is 1860

1-1 Child process is under execution

1-2 Process id of the child process is 1865

1-3 Process id of the parent process is 1861